

CS5542: Big Data Analytics and Apps

LAB Assignment 11

Show and Tell

- a deep neural network that learns how to describe the content of images
- an example of an encoder-decoder neural network
- first "encoding" an image into a fixed-length vector representation
- "decoding" the representation into a natural language description
- We will use inception v3 for image encoding
- The decoder is a long short-term memory (LSTM) network.
- LSTM network is trained as a language model conditioned on the image encoding
- Words in the captions are represented with an embedding model. Each word in the vocabulary is associated with a fixed-length vector representation that is learned during training.

First Phase of Training • parameters of the Inception v3 model are kept fixed

- it is simply a static image encoder function
- A single trainable layer is added on top of the Inception v3 model to transform the image embedding into the word embedding vector space
- The model is trained with respect to the parameters of the word embeddings, the parameters of the layer on top of Inception v3 and the parameters of the LSTM

The program is executed using the models generated from the first phase training on the images.

Here are the screenshots:









PyCharm Community Edition 2016.3.2 interface showing the 'ShowAndTell' project and the 'run_inference.py' script.

Project Structure:

- data
- inference_utils
- ops
- pretrained_model
 - 252600.jpg
 - 258391.3.jpg
 - 3164328039.jpg
 - 3165123595.jpg
 - build_data.py
 - configuration.py
 - evaluate.py
 - img_5_00022038.jpg
 - img_31.jpg
 - img_000089.jpg
 - img_699.jpg
 - inference_wrapper.py
 - nlTK

run_inference.py Code:

```
18 from __future__ import division
19 from __future__ import print_function
20
21 import math
22 import os
23
24
25 import tensorflow as tf
26
27 import configuration
28 import inference_wrapper
29 from inference_utils import caption_generator
30 from inference_utils import vocabulary
31
32 FLAGS = tf.flags.FLAGS
33
34 tf.flags.DEFINE_string("checkpoint_path", "pretrained_model/",
35                       "Model checkpoint file or directory containing a "
36                       "model checkpoint file.")
37
38 tf.flags.DEFINE_string("vocab_file", "pretrained_model/word counts.txt", "Text file containing the vocabulary.")
39 tf.flags.DEFINE_string("input_files", "258391.3.jpg,img_000089.jpg,252600.jpg,img_5_00022038.jpg",
40                       "File pattern or comma-separated list of file patterns "
41                       "of image files.")
```

Run Output:

Captions for image img_000089.jpg:

- 0) a pitcher in a blue and white uniform is throwing a baseball . (p=0.000499)
- 1) a pitcher in a blue and white uniform throwing a baseball . (p=0.000177)
- 2) a pitcher in a red and white uniform is throwing a baseball . (p=0.000086)
- 3) a pitcher in a red and white uniform throws a baseball . (p=0.000076)

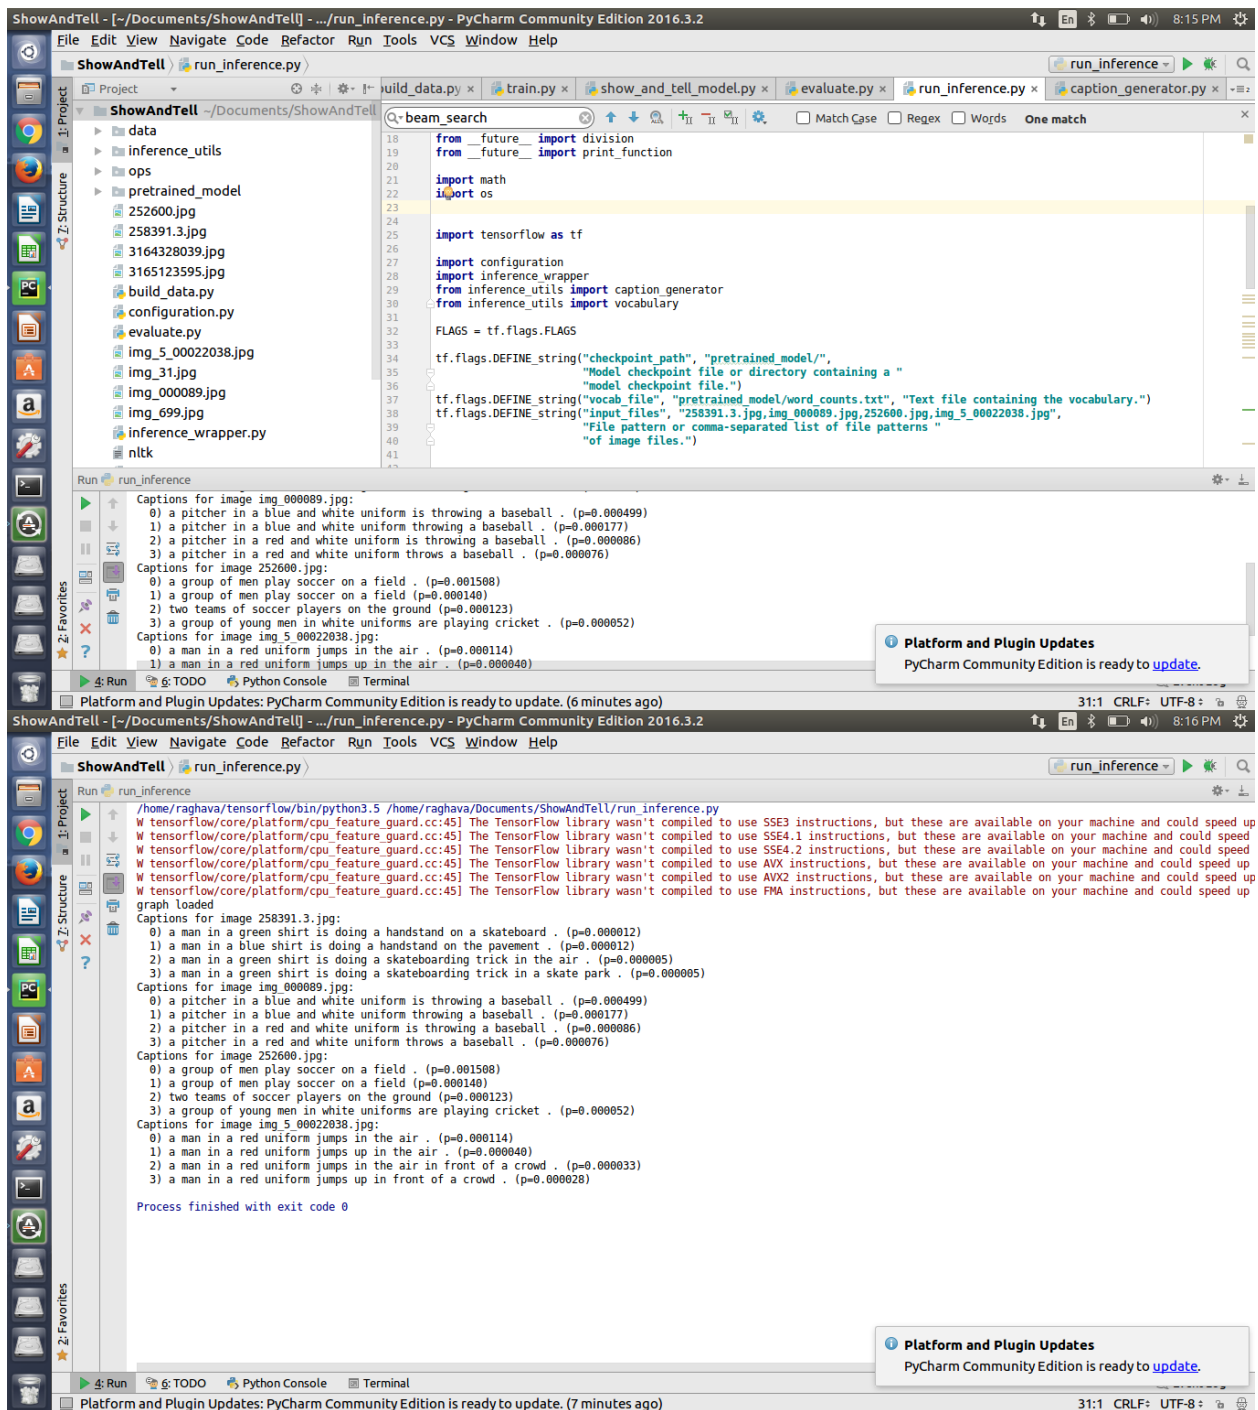
Captions for image 252600.jpg:

- 0) a group of men play soccer on a field . (p=0.001508)
- 1) a group of men play soccer on a field (p=0.000140)
- 2) two teams of soccer players on the ground (p=0.000123)
- 3) a group of young men in white uniforms are playing cricket . (p=0.000052)

Captions for image img_5_00022038.jpg:

- 0) a man in a red uniform jumps in the air . (p=0.000114)
- 1) a man in a red uniform jumps up in the air . (p=0.000040)

Platform and Plugin Updates: PyCharm Community Edition is ready to [update](#).



PyCharm Community Edition 2016.3.2 interface showing two screenshots of image captioning results.

Screenshot 1 (Top): The main window displays the image `252600.jpg` (900x600 JPEG, 24-bit color, 139.01K). The project structure on the left includes `ShowAndTell` with subfolders `data`, `inference_utils`, `ops`, and `pretrained_model`, and files `252600.jpg`, `258391.3.jpg`, `3164328039.jpg`, `3165123595.jpg`, `build_data.py`, `configuration.py`, `evaluate.py`, `img_5_00022038.jpg`, `img_31.jpg`, `img_000089.jpg`, `img_699.jpg`, `inference_wrapper.py`, and `nlTK`. The Run console shows the following captions for `image 252600.jpg`:

- 0) a group of men play soccer on a field . (p=0.001508)
- 1) a group of men play soccer on a field (p=0.000140)
- 2) two teams of soccer players on the ground (p=0.000123)
- 3) a group of young men in white uniforms are playing cricket . (p=0.000052)

The Run console also shows captions for `image img_5_00022038.jpg`:

- 0) a man in a red uniform jumps in the air . (p=0.000114)
- 1) a man in a red uniform jumps up in the air , (p=0.000040)
- 2) a man in a red uniform jumps in the air in front of a crowd . (p=0.000033)
- 3) a man in a red uniform jumps up in front of a crowd . (p=0.000028)

The process finished with exit code 0. A notification at the bottom right states: "Platform and Plugin Updates: PyCharm Community Edition is ready to update."

Screenshot 2 (Bottom): The main window displays the image `img_5_00022038.jpg` (720x240 JPEG, 24-bit color, 19.91K). The project structure on the left is identical to the first screenshot. The Run console shows the same captions for `image 252600.jpg` and `image img_5_00022038.jpg` as in the first screenshot. The process finished with exit code 0. A notification at the bottom right states: "Platform and Plugin Updates: PyCharm Community Edition is ready to update."

PyCharm Community Edition 2016.3.2 interface showing image captioning results for two images.

Image 1: 258391.3.jpg (900x506 JPEG, 146.81K)

Run `run_inference` results:

- graph loaded
- Captions for image 258391.3.jpg:
 - 0) a man in a green shirt is doing a handstand on a skateboard . (p=0.000012)
 - 1) a man in a blue shirt is doing a handstand on the pavement . (p=0.000012)
 - 2) a man in a green shirt is doing a skateboarding trick in the air . (p=0.000005)
 - 3) a man in a green shirt is doing a skateboarding trick in a skate park . (p=0.000005)
- Captions for image img_000089.jpg:
 - 0) a pitcher in a blue and white uniform is throwing a baseball . (p=0.000499)
 - 1) a pitcher in a blue and white uniform throwing a baseball . (p=0.000177)
 - 2) a pitcher in a red and white uniform is throwing a baseball . (p=0.000086)
 - 3) a pitcher in a red and white uniform throws a baseball . (p=0.000076)
- Captions for image 252600.jpg:
 - 0) a group of men play soccer on a field . (p=0.001508)

Image 2: img_000089.jpg (312x450 JPEG, 21.61K)

Run `run_inference` results:

- Captions for image img_000089.jpg:
 - 0) a pitcher in a blue and white uniform is throwing a baseball . (p=0.000499)
 - 1) a pitcher in a blue and white uniform throwing a baseball . (p=0.000177)
 - 2) a pitcher in a red and white uniform is throwing a baseball . (p=0.000086)
 - 3) a pitcher in a red and white uniform throws a baseball . (p=0.000076)
- Captions for image 252600.jpg:
 - 0) a group of men play soccer on a field . (p=0.001508)
 - 1) a group of men play soccer on a field (p=0.000140)
 - 2) two teams of soccer players on the ground (p=0.000123)
 - 3) a group of young men in white uniforms are playing cricket . (p=0.000052)
- Captions for image img_5_00022038.jpg:
 - 0) a man in a red uniform jumps in the air . (p=0.000114)
 - 1) a man in a red uniform jumps up in the air . (p=0.000040)

The accuracy turned out to be good enough, but not high enough because of the models are from only the first phase training. The above images generated fairly accurate captions.

Heroku Web Conversation-Tensor Flow

The second part of the program is Web Conversation using tensor flow, node js, api.ai and google conversation.

Here are the screenshots for the program:

The image shows a Heroku deployment terminal window and the API.AI Fulfillment console interface.

Terminal Output:

```
C:\Users\raghava koundinya>cd F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\ClientApp\target
C:\Users\raghava koundinya>F:
F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\ClientApp\target>heroku plugins:install heroku-cli-deploy
Installing plugin heroku-cli-deploy... done
F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\ClientApp\target>
F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\ClientApp\target>heroku deploy:jar service.war --app lab11a1
Uploading service.war
-----> Packaging application...
- app: lab11a1
- including: service.war
-----> Creating build...
- file: slug.tgz
- size: 2MB
-----> Uploading build...
- success
-----> Deploying...
remote: -----> heroku-deploy app detected
remote: -----> Installing OpenJDK 1.8... done
remote: -----> Discovering process types
remote: Procfile declares types -> web
remote: -----> Compressing...
remote: Done: 51M
remote: -----> Launching...
remote: https://lab11a1.herokuapp.com/ deployed to Heroku
remote: -----> Done
F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\ClientApp\target>
```

API.AI Fulfillment Console:

- Webhook:** Your web service will receive a POST request from API.AI in the form of the response to a user query matched by intents with webhook enabled. Be sure that your web service meets all the webhook requirements.
- Webhook example:** [Webhook example](#)
- ENABLED:** ☒
- URL:**
- BASIC AUTH:** Enter username... Enter password...
- HEADERS:** Enter key... Enter value...
[+ Add header](#)
- DOMAINS:** Disable webhook for all domains

Chat Interface:

- Intents:** lab11a
- Entities:** +
- Training [beta]**
- Integrations**
- Fulfillment**
- Prebuilt Agents**
- Account**
- Logout**

Chat Log:

- Svetlana from API.AI:** Hi Raghava, Great progress on building your agent! Now is a good time to tune some machine learning settings for better performance. Read more about them [here](#). Let me know if you have any questions.

Write a reply...

```
npm
-----> Creating build...
- including: service.war
- file: slug.tgz
- size: 2MB
-----> Uploading build...
- success
-----> Deploying...
remote: -----> heroku-deploy app detected
remote: -----> Installing OpenJDK 1.8... done
remote: -----> Discovering process types
remote: Procfile declares types -> web
remote: -----> Compressing...
remote: Done: 51M
remote: -----> Launching...
remote: Released v4
remote: https://lab11a1.herokuapp.com/ deployed to Heroku
remote: -----> Done

F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\ClientApp\target>cd react-app
The system cannot find the path specified.

F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\ClientApp\target>cd F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master
F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master>cd React-app
F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\React-App>npm install
[ ..... ] \ fetchMetadata: warn afterAdd C:\Users\raghava koundinya\AppData\Roaming\npm-cache\fs.realpath\1.0.0\package\package.json writtentenen
```

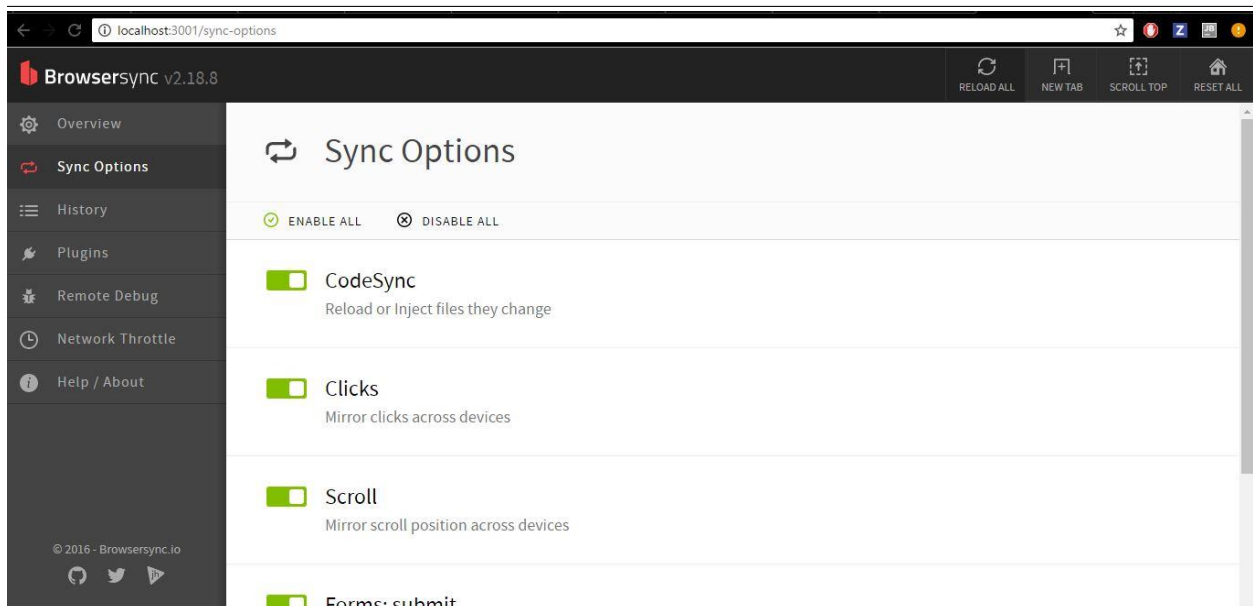
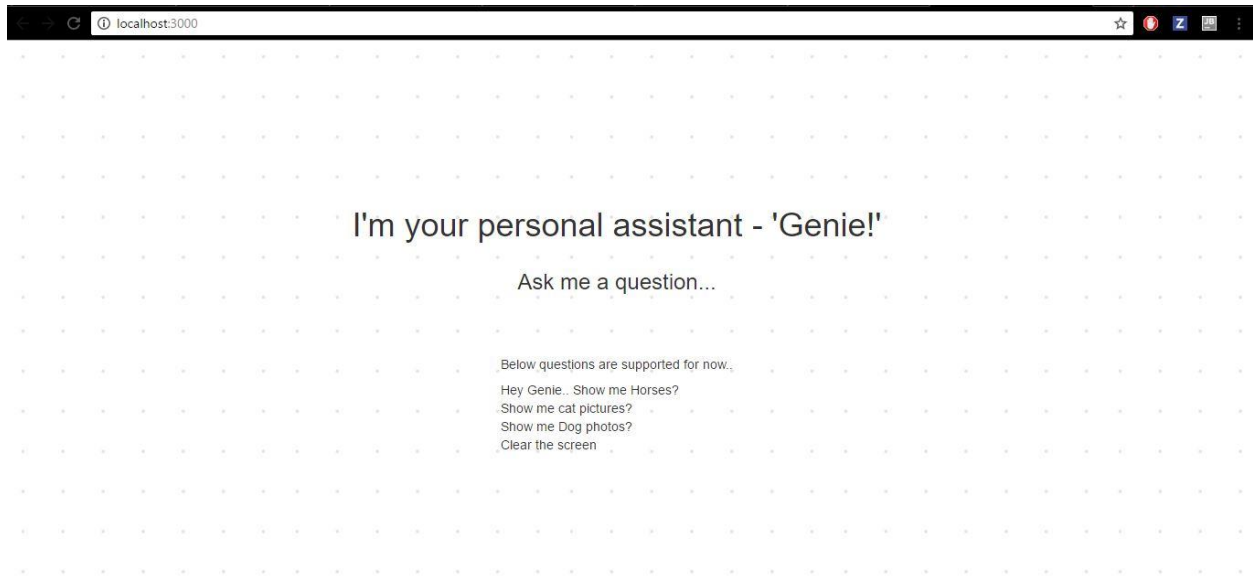
```
npm
| | -- vm-browserify@0.0.4
| +- supports-color@3.2.3
| | -- has-flag@1.0.0
| +- tapable@0.2.6
| +- watchpack@1.3.1
| | -- async@2.3.0
| | -- webpack-sources@0.2.3
| | -- source-list-map@1.1.1
+- webpack-dev-middleware@1.10.1
| +- mime@1.3.4
| +- range-parser@1.2.0
| +- webpack-hot-middleware@2.18.0
| +- ansi-html@0.7.7
| +- html-entities@1.2.0
| +- querystring@0.2.0
| -- strip-ansi@3.0.1

npm WARN optional SKIPPING OPTIONAL DEPENDENCY: fsevents@~1.0.0 (node_modules\chokidar\node_modules\fsevents):
npm WARN notsup SKIPPING OPTIONAL DEPENDENCY: Unsupported platform for fsevents@1.1.1: wanted {"os":"darwin","arch":"any"} (current: {"os":"win32","arch":"x64"})
npm WARN dynamic_client_view@1.0.0 No repository field.

F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\React-App>npm start
'npm' is not recognized as an internal or external command,
operable program or batch file.

F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\React-App>npm start

> dynamic_client_view@1.0.0 prestart F:\BigDataAnalytics\Tutorials\lab 11\Conversation-client-master\React-App
> webpack
```



Actions on Google

DesignDevelopDistributeSupport

Search

All Products

GUIDESSAMPLESREFERENCEWEB SIMULATOR

Please turn on Voice & Audio Activity for your Google Account

show me tulips

Please turn on Voice & Audio Activity for your Google Account

show me tulips

Please turn on Voice & Audio Activity for your Google Account

Type something, or click the mic and speak

```
{
  "query": "show me tulips",
  "accessToken": "ya29.G10s80T7Iqb70cFuIw9nn5pSTbhQ-n-
MwKDvcBUfrUib_tISNuMR0KnBE0VA0Vn6nOv40r4UI0xJ9vpLh8b6aI9cHmp9xscooofF4dysuGD7VVhnPN
RY70qefMDH8hU"
}
```

Response

```
{
  "response": "Please turn on Voice & Audio Activity for your Google Account",
  "audioResponse": "//NEXAAQq...",
  "debugInfo": {
    "sharedDebugInfo": [
      {
        "name": "PermissionFailure",
        "debugInfo": "Please make sure your activities like Voice & Audio
Activity are enabled in https://support.google.com/websearch/answer/6030020?hl=en"
      }
    ]
  }
}
```

<https://developers.google.com/actions/>